

# Implement Gradient Descent in Python

## What is gradient descent ?

It is an optimization algorithm to find the minimum of a function. We start with a random point on the function and move in the **negative direction** of the **gradient of the function** to reach the **local/global minima**.

## Example by hand :

**Question :** Find the local minima of the function  $y=(x+3)^2$  starting from the point  $x=2$

**Solution :** We know the answer just by looking at the graph.  $y = (x+3)^2$  reaches it's minimum value when  $x = -3$  (i.e when  $x=-3, y=0$ ). Hence  $x=-3$  is the local and global minima of the function.

Now, let's see how to obtain the same numerically using gradient descent.

**Step 1 :** Initialize  $x = 2$ . Then, find the gradient of the function,  $dy/dx = 2*(x+3)$ .

**Step 2 :** Move in the direction of the negative of the gradient. Stopping condition: how much to move? For that, we require a learning rate. Let us assume the **learning rate**  $\rightarrow 0.01$

**Step 3 :** Let's perform 2 iterations of gradient descent

Initialize Parameters:

$$x_0 = 2$$

$$\text{Learning rate} = 0.01$$

$$\frac{dy}{dx} = \frac{d}{dx}(x+3)^2 = 2 \times (x+3)$$

Iteration 1:

$$x_1 = x_0 - (\text{learning rate}) \times \left(\frac{dy}{dx}\right)$$

$$x_1 = 2 - (0.01) \times (2 \times (2 + 3)) = 1.9$$

Iteration 2:

$$x_2 = x_1 - (\text{learning rate}) \times \left(\frac{dy}{dx}\right)$$

$$x_2 = 1.9 - (0.01) \times (2 \times (1.9 + 3)) = 1.802$$

**Step 4 :** We can observe that the X value is slowly decreasing and should converge to -3 (the local minima). However, how many iterations should we perform?

Let us set a precision variable in our algorithm which calculates the difference between two consecutive “x” values . If the difference between x values from 2 consecutive iterations is lesser than the precision we set, stop the algorithm !

## Gradient descent in Python :

**Step 1 :** Initialize parameters

```
cur_x = 2 # The algorithm starts at x=2
rate = 0.01 # Learning rate
precision = 0.000001 #This tells us when to stop the algorithm
previous_step_size = 1 #
max_iters = 10000 # maximum number of iterations
iters = 0 #iteration counter
df = lambda x: 2*(x+3) #Gradient of our function
```

**Step 2 :** Run a loop to perform gradient descent :

i. Stop loop when difference between x values from 2 consecutive iterations is less than 0.000001 or when number of iterations exceeds 10,000

```
while previous_step_size > precision and iters < max_iters:
    prev_x = cur_x #Store current x value in prev_x
    cur_x = cur_x - rate * df(prev_x) #Grad descent
    previous_step_size = abs(cur_x - prev_x) #Change in x
    iters = iters+1 #iteration count
    print("Iteration",iters,"\nX value is",cur_x) #Print iterations

print("The local minimum occurs at", cur_x)
```

**Output :** From the output below, we can observe the x values for the first 10 iterations- which can be cross checked with our calculation above. The algorithm runs for **571** iterations before it terminates. The code and solution is embedded below for reference.

```
X value is -2.9999364493143186
Iteration 559
X value is -2.9999377203280324
Iteration 560
X value is -2.999938965921472
Iteration 561
X value is -2.9999401866030424
Iteration 562
X value is -2.9999413828709818
Iteration 563
X value is -2.999942555213562
Iteration 564
X value is -2.999943704109291
Iteration 565
X value is -2.999944830027105
Iteration 566
X value is -2.999945933426563
Iteration 567
X value is -2.999947014758032
```

## LP-III, ML, Experiment 4: Implement Gradient Descent Algorithm

Iteration 568

X value is -2.9999480744628713

Iteration 569

X value is -2.999949112973614

Iteration 570

X value is -2.999950130714142

Iteration 571

X value is -2.999951128099859

The local minimum occurs at -2.999951128099859

### Screenshot of program and output

The screenshot displays the Spyder Python IDE interface. The editor window shows a Python script named `gd.py` implementing the Gradient Descent algorithm. The script includes comments for encoding, creation date, and author. The algorithm starts with `cur_x = 2`, a learning rate of `0.01`, and a precision threshold of `0.000001`. It iterates up to `10000` times, updating `cur_x` based on the gradient `df`. The output shows iterations from 562 to 571, with the final value of `cur_x` being `-2.999951128099859`.

```
1  #-*- coding: utf-8 -*-
2  """
3  Created on Wed Sep 21 11:19:44 2022
4
5  @author: dell: Prof Vijay More
6  """
7
8  cur_x = 2 # The algorithm starts at x=2
9  rate = 0.01 # Learning rate
10 precision = 0.000001 #This tells us when to stop the algorithm
11 previous_step_size = 1 #
12 max_iters = 10000 # maximum number of iterations
13 iters = 0 #iteration counter
14 df = lambda x: 2*(x+3) #Gradient of our function
15
16 while previous_step_size > precision and iters < max_iters:
17     prev_x = cur_x #Store current x value in prev_x
18     cur_x = cur_x - rate * df(prev_x) #Grad descent
19     previous_step_size = abs(cur_x - prev_x) #Change in x
20     iters = iters+1 #iteration count
21     print("Iteration",iters,"\nX value is",cur_x) #Print iterations
22
23 print("The local minimum occurs at", cur_x)
24
```

The Variable Explorer on the right shows the current state of variables:

Name	Type	Size	Value
cur_x	float	1	-2.9999511280998...
iters	int	1	571

The Console window shows the output of the program:

```
X value is -2.9999401866030424
Iteration 562
X value is -2.9999413828709818
Iteration 563
X value is -2.999942555213562
Iteration 564
X value is -2.999943704109291
Iteration 565
X value is -2.999944830027105
Iteration 566
X value is -2.999945933426563
Iteration 567
X value is -2.999947014758032
Iteration 568
X value is -2.9999480744628713
Iteration 569
X value is -2.999949112973614
Iteration 570
X value is -2.999950130714142
Iteration 571
X value is -2.999951128099859
The local minimum occurs at -2.999951128099859
```

Source: <https://towardsdatascience.com/implement-gradient-descent-in-python-9b93ed7108d1>